ECE6703J

Computer-Aided Design of Integrated Circuits

Maze Routing Implementation Issues

Outline

- Maze Routing Implementation Issues
  - Expansion Step
  - Data Structures
  - Depth-First Search
  - Global Routing

Implementation Concerns

- Algorithms
  - Question: if we can only process one cell at a time…
  - …then, which cell is next to "label" in the search process

- Data structure (representation)
  - How do we store the routing grid?
  - What do we need in each cell?
  - How do we represent the state of the path search process?

First Issue: Expansion Step

- One big goal
  - Efficient storage: Big layout needs a big grid; put as little in each cell as possible.

- Big Idea: **Do not** store path costs in grid cells
  - If we have big path cost, then we need many bits per cell.
  - <u>Notice</u>: Only cells **most recently labeled** during search are used to **expand** the search.
  - Important terminology: these cells comprise the **search wavefront**.
    - It is the **wavefront** that is really most important....

# Important Idea: Search Wavefront

- What is wavefront?
  - Set of cells already **reached** in the expansion process…
  - …that **have neighbors** that we have **not** yet reached.
  - Wavefront is the **frontier** of the active search for new paths.
- Implication
  - **Don't** store the pathcost numbers **in the grid**.
  - Just store the wavefront cells themselves in a **separate data structure**.

|	3	2	3		
3	2	1	2	3	
2	1	S	1	2	3
3	2	1	2	3	
	3	2	3		
		3			

How to Expand Wavefront?

Answer: Expand in the **pathcost** order

After expanding neighbors of the cost "**0**" grid, i.e., S


After expanding neighbors of the cost "**1**" grid

cost 1   cost 3

More Complex Wavefront

After expanding neighbors of the cost "**2**" grid

	3	2	3		
3	2	1	2	3	
2	1	S	3		
3	2	1	4		
	3	2	3		
		3			

☐ cost 1   ▨ cost 3

- **Expanded** in **pathcost order**, cheapest cells before more expensive cells.
  - Algorithm: **Cheapest-cell-first search** (a variants of Dijkstra's algorithm)

Cheapest-Cell-First Search

- Find a **cheapest** cell C from the **wavefront**.

- Find the neighbors $\{N1, N2, \ldots Nk\}$ of cell C that you have **not reached** yet.

- Compute the cost of expanding through C to reach the cells in $\{N1, N2, \ldots Nk\}$.

- Add new cells $\{N1, N2, \ldots Nk\}$ to the wavefront data structure, including their newly computed pathcosts, $\{pathcost(N1), pathcost(N2), \ldots, pathcost(Nk)\}$.
  - Also, add to each of these new cells a pointer to cell C as the **predecessor**.
  - Mark the routing grid to remember that we have **reached** these cells $\{N1, N2, \ldots Nk\}$.

- Remove cell C from the wavefront.

- Repeat with the **next cheapest cell on wavefront**…

Cell Predecessor

- **Predecessor** of cell A on the wavefront is a "tag" that tells **direction** from which cell A was reached, i.e., it's how we found A.

- Why do we have to remember this?
  - Because we **do not mark pathcosts** in the grid any longer.
  - Then, to backtrace the path, we cannot simply follow the numbers in decreasing order.
  - We need a real "trail" that points, cell to cell, from target back to source.

Marking the Predecessor

- Assume 2 routing layers. 6 pred(A) tags:
  N, S, E, W, Up, Down
  - We only need 3 bits to store 6 tags.
  - **<u>Note</u>**: pred(A) is the direction from which cell A was reached

expand → reach

B
A

prec(A)=N

reach ↑ expand

A
B

prec(A)=S

expand ←
reach

A	B

prec(A)=E

expand →
reach

B	A

prec(A)=W

Marking the Predecessor (cont.)

layer 2   B   expand

↓

layer 1   A   reach

prec(A)=Up

layer 2   A   reach

↑

layer 1   B   expand

prec(A)=Down

We Store Predecessors on Grid

Layer 1

Layer 2

- Result: backtrace is easy! Follow pred(C) tags in the grid.

Example: Expanding & Reaching

B: next to expand
Know:
pathcost(B) = 12
pred(B)=N

A: unreached
Know:
cell cost=3
Compute:
pathcost(A)
pred(A)

- Expand cell B, reach cell A
  - B is with **cheapest** cost on wavefront.
  - A is B's **unreached** neighbor.
  - Compute cost to reach A from B: 12+3=15
  - Add this new "cell object" to the wavefront:

  $<$layer=L; cell location=(x_A, y_A); pathcost=15$>$

  - Mark cell A as Reached(A)=true
    - In future, we won't try to put it on the wavefront **again.**
  - In cell A in the gird, set pred(A) = W

Basic Maze Routing Algorithm

wavefront = { source cells };
while (we have not reached target cell) {
    if (**wavefront** == empty) quit; *// no path to be found*
    C = **get** lowest cost cell on **wavefront** structure;
    if (**C** == target) {
        backtrace path in grid; *// follow pred( ) pointers to source*
        cleanup & **return**; *// we found a path*
    }
    foreach (**unreached** neighbor **N** of cell **C**) {
        mark cell **N** in grid as **reached**.
        mark **pred(N)** with direction from cell **N** to **C**;
        // compute cost to reach
        pathcost(N) = pathcost(C) + cellcost(N) ;
        add cell **N** to wavefront, indexed by **pathcost(N)**;
    }
    delete cell **C** from wavefront;
}

Observations

- Core algorithm is fairly simple. Works for grid with non uniform costs.

- Extends naturally to other features
  - **Multi-point nets**: just mark all the cells of intermediate paths as source, put them all in the wavefront to route to next point.
  - **Multiple routing layers**: just use parallel grids. When you expand to reach neighbors, check UP, DOWN to other layers. Use pred(Cell) = UP, DOWN to do backtrace.

Outline

- Maze Routing Implementation Issues
  - Expansion Step
  - Data Structures
  - Depth-First Search
 - Global Routing

16

Two Key Data Structures: Overview

- Routing grid
  - Routing surface. Contains blockage. Has costs of each cell.
  - Mark these cells if they have already been reached.
  - Mark predecessor.

- Wavefront
  - Holds active cells **to be expanded**.
  - Cells store pathcost.
  - Indexed on pathcost; always expand cell with cheapest pathcost next.

Routing Grid

- Each grid cell C stores:
  - (x, y) coordinate of cell C
  - Layer of grid cell C
  - Cost(C) (a small number)
    - **Note**: we can use cost $= -1$ to indicate blockage
  - Pred(C) tag $= N,S,E,W,Up,Down$
  - Reached(C), 1 bit Boolean true/false

Wavefront

- Each cell C in wavefront stores:
  - (x, y) coordinate of cell C
  - Layer of grid cell C
  - Pathcost(C) = length of the shortest path to cell C

- We need something clever here:
  - Need fast **extracting minimum** operations, to take a cell with the smallest pathcost from wavefront
  - Need fast insertion, to add a cell into the wavefront

Wavefront Structure: Min Heap

- Store cells of wavefront in a **min heap**.

  - Also called a **min priority queue**.

  - Classical data structure designed for fast insertion and extracting minimum operations.

    - These two operations have $O(\log N)$ time complexity for N objects.

Recall: Min Heap

- A min heap is a **complete binary tree** that
  - For **any** node v, the key of v is smaller than or equal to ($\leq$) the keys of any **descendants** of v.

- The item with the smallest cost is the **root**!

- Insert operation "bubbles up" the new item in heap to ensure the min heap property.

- Extracting minimum operation also ensures the min heap property.

Outline

- Maze Routing Implementation Issues
  - Expansion Step
  - Data Structures
  - Depth-First Search
  - Global Routing

22

Expansion Process, Revisited

- Problem:
  - Expand lots of cells to find one path to the target.
  - CPU time is proportional to number of cells you expand.
  - No attempt to search in direction of target first.

- Questions:
  - Can we actually bias expansion so we search **toward** the target?
  - Can we do this and still keep **guarantees** of reaching target with minimum cost path?

Motivation for Smart Search

Expanding **away** from the target seems to be a waste of time.

Searching **toward** the target in the shaded region, sometimes called the "source-target box", seems a good idea to try first.

Smarter Search: Maze Search With Predictor

- Add **predictor function** to the cost, to direct the search **toward** the target.

- Plain maze router

  - Add a cell C to wavefront with cost that measures cost of **partial path**, source-to-cell(C).

- Smarter maze router

  - Add cell C to wavefront with cost that **estimates entire source-to-target cost** of path **through the cell**.

  - Trick: estimate this as
    pathcost(source to cell C) **+ predictor(cell C to target)**.

  - We call this method as smart depth first search (DFS) with a predictor in contrast to classical breadth first search (BFS).

# Plain Maze Routing

Source

Expand

Reach

Target

Reached cell goes on wavefront with this cost:

pathcost(source to cell being expanded) + cost of newly reached cell

Add A Depth-First Predictor

Source

Expand

Reach

Target

Reached cell goes on wavefront with this cost:

pathcost(source to cell being expanded) + cost of newly reached cell

+ estimate of pathcost(from newly reached cell to target)

A typical estimator is:
(Manhattan distance to target) × (Smallest cell cost)

Technical Results

- Depth first predictor
  - It is famous application of a classical idea: **A* search**.
  - Since predictor is a **lower bound** on extra pathcost you add to reach target…
  - …by the property of A* search, you are **guaranteed** to get the minimum cost path.

- What does it do?
  - It **alters the order** in which we expand cells.
 - It prefers to expand cells that are **closer to the target first**.

Depth-First Expansion Example

Pathcost = 1;
Distance to target = 5
Esitmate of cost: 1+5=6

Observation: Search prefers to stay **inside** the bounding box of the source-target rectangle before it expands other cells.

Pathcost = 1;
Distance to target = 3
Esitmate of cost: 1+3=4

Depth-First Expansion Example

- It turns out for this example all the cells in this box cost the same!

  - **In this box**:

  pathcost(source→cell) + estimated_pathcost(cell→target) = constant

  - So – it expands **toward** the target!

|					
		S	4	4	
		4	4	4	
		4	4	T	

DFS Expansion: Unintended Consequences

- This is "the" example of which "stay in the box" is **inefficient**.
  - The target is blocked inside the source-target box.
  - **Problem**: DFS explores the **whole shaded** rectangle before it tries any place else.
  - In this case, it may be faster to search outside the box!
- But, generally, this heuristic is good!

A Variation of the Heuristic

Source

Expand

Reach

Target

Reached cell goes on wavefront with this cost:

pathcost(source to cell being expanded) + cost of newly reached cell

$+\ \boxed{\mathbf{K} \times}\ $ estimate of pathcost(from newly reached cell to target)

K is usually (1 + a very small number)
<u>Idea</u>: Insist that expand in a cheaper path toward target. Distorts optimum a bit, but it is faster.

Outline

- Maze Routing Implementation Issues
  - Expansion Step
  - Data Structures
  - Depth-First Search
 - Global Routing

33

Reality: ASIC Scale & Complexity

- Big chip, 1cm x 1cm

- 20-50 million nets

- Modern IC technology, ~100nm pitch for wires (grid size)

- So, 100K x 100K routing grid!

- 10 routing layers

- 100 billion grid cells!

- Do we really do it like this?

No!

One Way: Preferred Routing Directions

- Every layer of metal wiring has a **preferred direction**.
  - E.g., metal 3, 5, 7, 9 prefer **vertical**; metal 4, 6, 8, 10 prefer **horizontal**.
  - As a result, all the wires on these layers look like straight lines.
  - If you need to bend, use a **via**. Makes it easier to embed lots of wires.
- This is **<u>one way</u>** to reduce complexity.

Divide & Conquer: Global Routing

- All the routing we have done so far is called **Detailed Routing**, which plans **exact, final** path of wire.

- How to deal with huge scale of chips? **Global Routing**
  - Global routing also imposes a grid on surface of the chip.
  - But now, this grid is much coarser, e.g., each grid cell is about 200 wire grids × 200 wire grids in size.

- Big Idea: **Maze route** through these big, coarse regions.

Global vs Detailed Routing: Geometry

- Each cell in **Global Routing** grid is a box (called a **GBOX**) with size typically $100{\sim}200$ wire grids on a side.
  - This example just has 6 for clarity.

What Global Routing Does?

- Global routing tells us we want this net to use this **rough** path; but not on which **exact** path inside this region.
  - Generate **regions of confinement** (i.e., coarse path) for a wire.
- Makes sure overall **wiring congestion** is reasonable.
  - Balance **supply** (how much available space) vs **demand** (how many paths want to go here).

What Detailed Routing Does?

- Detailed routing embeds **exact** paths in these regions.
 - Simplest model: **a finer grid**. Require wires to use tracks on this grid.

This reduces computation!

Global Routing as Maze Routing

- What's different for a global router? **Cell Cost Function**
  - **Detailed router**: grid cells are blocked (cost=infinity) or they have a cost proportional to how much we want to avoid them.
  - **Global router**: grid cells have **dynamic** cost, proportional to how many **previously** routed wires used this cell, how many more can fit **later**.

Suppose global GBOX is 10 tracks for wires.

GBOX Cost

∞

low

0 1 2 3 4 5 6 7 8 9 10 11 12   #previous wires

Maze Routing: Summary

- Has been around a long time.
  - Very **flexible** cost-based search; can be recast to attack many problems.
  - **Dominates** most modern routers…
  - Some other approaches, e.g., sophisticated "grid-less" representation of "space".
- Lots of ancillary problems (and solutions)
  - Hierarchy/abstraction (e.g., global routing).
  - Iterative improvement strategies.